


# Training neural network

# Gradient descent


- ▶ Hill climbing algorithm that follows the gradient of the function to be optimized



see notebook Linear - Nonlinear separation

# Gradient descent

- ▶ Hill climbing algorithm that follows the gradient of the function to be optimized



$\mathbf{w} \leftarrow$  any point in the parameter space

**loop** until convergence **do**

**for each**  $w_i$  **in**  $\mathbf{w}$  **do**



$$w_i \leftarrow w_i - \alpha \frac{\partial}{\partial w_i} \text{Loss}(\mathbf{w})$$

see notebook Linear - Nonlinear separation

# Linear regression - with gradient descent

- ▶ Univariate linear function with input  $x$  and output  $y$ :  
 $y = w_1 * x + w_0$
- ▶ Univariate Linear regression - finding the  $h_w$  that best fits  $n$  point  $x, y$

$$h_w(x) = w_1 x + w_0$$



How? with gradient descent in order to minimize the loss  
[see detailed analysis](#) - it will also be discussed during IS lab

# Math explanation

$$Loss(h_{\mathbf{w}}) = \sum_{j=1}^N L_2(y_j, h_{\mathbf{w}}(x_j)) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(x_j))^2 = \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2$$

For one example:

$$\begin{aligned}\frac{\partial}{\partial w_i} Loss(\mathbf{w}) &= \frac{\partial}{\partial w_i} (y - h_{\mathbf{w}}(x))^2 \\ &= 2(y - h_{\mathbf{w}}(x)) \times \frac{\partial}{\partial w_i} (y - h_{\mathbf{w}}(x)) \\ &= 2(y - h_{\mathbf{w}}(x)) \times \frac{\partial}{\partial w_i} (y - (w_1 x + w_0))\end{aligned}$$

$$\frac{\partial}{\partial w_0} Loss(\mathbf{w}) = -2(y - h_{\mathbf{w}}(x)); \quad \frac{\partial}{\partial w_1} Loss(\mathbf{w}) = -2(y - h_{\mathbf{w}}(x)) \times x$$

Gradient descent rule:

$$w_0 \leftarrow w_0 + \alpha (y - h_{\mathbf{w}}(x)); \quad w_1 \leftarrow w_1 + \alpha (y - h_{\mathbf{w}}(x)) \times x$$


Batch gradient descent

$$w_0 \leftarrow w_0 + \alpha \sum_j (y_j - h_{\mathbf{w}}(x_j)); \quad w_1 \leftarrow w_1 + \alpha \sum_j (y_j - h_{\mathbf{w}}(x_j)) \times x_j$$

# Backpropagation

see slide 1 - what is a neural network?

# Neural networks



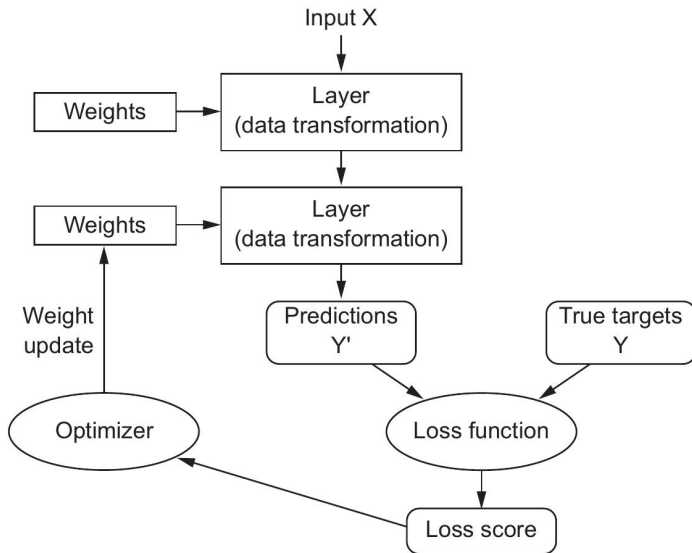
$$a_j = g(in_j)$$

$$a_j = g\left(\sum_i (w_{i,j} * a_i)\right)$$

- ▶  $a_i$  output of neuron  $i$ ,
- ▶  $a_0 = 1$  bias
- ▶  $W_{i,j}$  weight between neuron  $i$  and neuron  $j$
- ▶  $g$  activation function

example:  $a_{h_1} = g(w_{1,1} * a_i + w_{2,1} * a_2 + w_{0,1})$

# Layered transformation





# Loss


- Mean square error MSE

$$MSE = \frac{1}{N} \sum_{(x,y) \in D} (y - \text{prediction}(x))^2$$

- Log Loss

$$\text{LogLoss} = \sum_{(x,y) \in D} -y \log(y') - (1 - y) \log(1 - y')$$

- $(x, y) \in D$  - labeled dataset;  $y$  - label for sample  $x$ ,  $x$  is a vector of features. For binary classification  $y$  is 0 or 1;  $y'$  -  $\text{prediction}(x)$



# Backpropagation - based on Gradient Descent

## Back-propagation

- ▶ Output layer

$$w_{j,i} = w_{j,i} + \alpha * a_j * \Delta_i$$

, unde  $\Delta_i = Error_i * g'(in_i)$

- ▶ Hidden layer ( $k$  node -  $j$  node -  $i$  nodes)  
Propagated error:

$$\Delta_j = g'(in_j) \sum_i w_{j,i} \Delta_i$$

- ▶ hidden node  $j$  is responsible for some fraction of the error  $\Delta_i$  in each of the output nodes to which it connects; stronger connection (higher weight) means higher contribution to error

$$w_{k,j} = w_{k,j} + \alpha * a_k * \Delta_j$$

# Backpropagation

- ▶ compute the  $\Delta$  values for the output units using the observed error (loss)
- ▶ starting with output layer, repeat for each layer in the network, until earliest hidden layer is reached:
  - ▶ propagate the *Delta* value back to the previous layer
  - ▶ update the weights between two layers

Math demo ( Chain rule (nice explanations AIMA, [notes](#))

# Neural Network Playground

<https://playground.tensorflow.org/>

# What can go wrong

- ▶ Vanishing gradient
- ▶ Exploding gradient
- ▶ Overfitting/underfitting
  - ▶ Adjust size of the network
  - ▶ Use regularization: reduce complexity

$$\text{minimize}(\text{Loss}(\text{Data}|\text{Model}) + \text{complexity}(\text{Model}))$$

- ▶ L2 penalizes  $\text{weight}^2$ , L1 penalizes  $||\text{weight}||$ .
  - ▶ the derivative of L2 is  $2 * \text{weight}$ , The derivative of L1 is  $k$  (a constant, whose value is independent of weight). L2 does not normally drive weights to zero, while L1 can  $\rightarrow$  L1 eliminates irrelevant features
- ▶ Dropout - Dropout, is useful for neural networks. It works by randomly "dropping out" unit activations in a network for a single gradient step. The more you drop out, the stronger the regularization:
- ▶ Early stopping
- ▶ Adjust learning rate

# Vanishing gradient

- ▶ The gradients for the lower layers (closer to the input) can become very small. In deep networks, computing these gradients can involve taking the product of many small terms.
- ▶ When the gradients vanish toward 0 for the lower layers, these layers train very slowly, or not at all.
- ▶ The ReLU activation function can help prevent vanishing gradients.

see Google Machine Learning Crash Course

# Exploding gradient

- ▶ If the weights in a network are very large, then the gradients for the lower layers involve products of many large terms. In this case you can have exploding gradients: gradients that get too large to converge.
- ▶ Batch normalization can help prevent exploding gradients, as can lowering the learning rate.